

TECHNISCHE UNIVERSITÄT DRESDEN

DEPARTMENT OF COMPUTER SCIENCE
INSTITUTE OF COMPUTER ENGINEERING
CHAIR OF COMPUTER ARCHITECTURE
PROF. DR. WOLFGANG E. NAGEL

Hauptseminar
“Rechnerarchitektur und Programmierung”

NCCL Profiler Plugin API – A Feasibility Study

Alexander Moritz Van Le
(Mat.-No.: 4607469)

Professor: Prof. Dr. Wolfgang E. Nagel
Tutor: Bert Wesarg

Dresden, February 25, 2026

Contents

| | |
|---|-----------|
| 1 Abstract | 2 |
| 2 Introduction to NCCL | 2 |
| 2.1 Comparison to MPI | 2 |
| 2.2 Relevant NCCL internals | 3 |
| 3 Profiler Plugin | 3 |
| 3.1 Profiler plugin loading mechanism | 6 |
| 3.2 Profiler API | 7 |
| 3.2.1 init | 8 |
| 3.2.2 startEvent | 9 |
| 3.2.3 stopEvent | 10 |
| 3.2.4 recordEventState | 13 |
| 3.2.5 finalize | 13 |
| 3.2.6 name | 14 |
| 4 Code examples and visualizations | 14 |
| 4.1 One Device per Thread | 16 |
| 4.2 Multiple Devices per Thread (ncclGroup) | 18 |
| 4.3 Aggregated operations | 20 |
| 5 Performance and scalability | 21 |
| 6 Discussion | 23 |
| 6.1 Considerations for developers | 23 |
| 6.2 Known limitations | 25 |
| 6.3 Potential Integration with Score-P | 25 |
| 7 Conclusion | 26 |
| References | 28 |

1 Abstract

2 Artificial intelligence (AI) has established itself as a primary use case in high-performance computing
3 (HPC) environments due to its compute-intensive and resource-intensive workloads. Analyzing and
4 optimizing application performance is therefore essential to maximize efficiency and reduce costs. Many
5 AI workloads involve communication between GPUs, often distributed across numerous GPUs in multi-
6 node systems. The NVIDIA Collective Communication Library (NCCL) serves as the core library for
7 implementing optimized communication primitives on NVIDIA GPUs. To provide detailed performance
8 insights, NCCL offers a flexible profiler plugin API. This allows developers to directly integrate custom
9 profiling tools into the library to extract detailed performance data on communication operations. This
10 feasibility study explores the capabilities and integration mechanisms of the API.

11 First, this study provides background information on NCCL, followed by an explanation of the Profiler
12 API accompanied with code examples and visualizations. Next, considerations for developers of the
13 Profiler API and its potential integration with Score-P is discussed. Finally, the study concludes with a
14 summary of the findings.

15 2 Introduction to NCCL

16 NCCL was first introduced by NVIDIA in 2015 at the Supercomputing Conference¹ with the code
17 published on GitHub[2]. The release of NCCL 2.0 in 2017 brought support for NVLink, however it was
18 initially only available as pre-built binaries. With the release of 2.3 in 2018, NCCL returned to being
19 fully open source on GitHub. The NCCL Profiler Plugin API was even later introduced with NCCL 2.23
20 in early 2025.

21 Before taking a closer look at the Profiler Plugin API, it is helpful to have some rudimentary understand-
22 ing on certain designs in NCCL.

23 2.1 Comparison to MPI

24 Although NCCL is inspired by the Message Passing Interface (MPI) in terms of API design and usage
25 patterns, there are notable differences due to their respective focuses:

- 26 • **MPI:** Communication is CPU-based. A rank corresponds to a single CPU process within a com-
27 municator.
- 28 • **NCCL:** Communication is GPU-based, with CPU threads handling orchestration. A rank corre-
29 sponds to a GPU device within a communicator, where the mapping from ranks to devices is sur-
30 jective. A single CPU thread can manage multiple ranks (i.e. multiple devices) in a communicator
31 using the functions `ncclGroupStart` and `ncclGroupEnd`. A CPU thread can also man-
32 age multiple ranks from *different* communicators (i.e. same device allotted by ranks from different

¹<https://images.nvidia.com/events/sc15/pdfs/NCCL-Woolley.pdf>

33 communicators through communicator creation with `ncclCommSplit` or `ncclCommShrink`),
 34 so the mapping from ranks to threads is also surjective.

35 **2.2 Relevant NCCL internals**

36 It helps to understand what NCCL does internally when an application calls the NCCL User API.

37 A typical NCCL application follows this basic structure:

- 38 • create nccl communicators
- 39 • allocate memory for computation and communication
- 40 • do computation and communication
- 41 • clean up nccl communicators

42 During NCCL communicator creation, NCCL internally spawns a thread called `ProxyService`. This
 43 thread lazily starts another thread called `ProxyProgress`², which handles network requests for GPU
 44 communication during collective and P2P operations. See Fig. 1.

45 if-guards ensure that these threads are created once per `ncclSharedResources`³. By default ev-
 46 ery NCCL communicator has its own shared resource. When the application calls `ncclCommSplit`
 47 or `ncclCommShrink`, where the original communicator was initialized with a
 48 `ncclConfig_t` with fields `splitShare` or `shrinkShare` set to 1, the newly created communi-
 49 cator uses the same shared resource (and the proxy threads) as the parent communicator.

50 Later, whenever the application calls the NCCL User API, NCCL internally decides what network op-
 51 erations to perform and calls `ncclProxyPost` to post them to a `proxyOpsPool`. See Fig. 2.

52 The `ProxyProgress` thread reads from this pool when calling `ncclProxyGetPostedOps` and
 53 progresses the ops. See Fig. 3.

54 Familiarity with this network activity pattern will aid in understanding the Profiler Plugin API’s behavior
 55 discussed in the following section.

56 **3 Profiler Plugin**

57 Whenever a communicator is created, NCCL looks for the existence of a profiler plugin and loads it
 58 if it has not already been loaded on the process. NCCL then initializes the plugin with the created
 59 communicator. Whenever the user application calls the Collectives or P2P API (e.g. `ncclAllReduce`)
 60 with this communicator, NCCL calls the profiler API in different regions of the internal code. When the

²`src/proxy.cc`.

³`src/include/comm.h`.




Figure 1: Thread creation: User API → NCCL internal init → create ProxyService → create ProxyProgress.




Figure 2: Flow from User API to ncclProxyPost




Figure 3: The ProxyThread runs in a simple progressing loop: progress ops, then get posted ops (or wait).

61 communicator is destroyed, the profiler plugin is unloaded if this was the only communicator on the
 62 process.

63 3.1 Profiler plugin loading mechanism

64 Each time a NCCL communicator is created, `ncclProfilerPluginLoad`⁴ is called, where NCCL
 65 looks for a shared library that represents the profiler plugin by checking an environment variable. It then
 66 calls `dlopen`⁵ and `dlSym` to load the library immediately with local symbol visibility:

```
67
68 profilerName = ncclGetEnv("NCCL_PROFILER_PLUGIN");
69 // ...
70 handle* = dlopen(name, RTLD_NOW | RTLD_LOCAL);
71 // ...
72 ncclProfiler_v5 = (ncclProfiler_v5_t*)dlsym(handle, "ncclProfiler_v5");
```

74 If the library has already been loaded on the process, this procedure is skipped.
 75 A `profilerPluginRefCount` keeps track of the number of calls to this procedure to ensure correct
 76 unloading during finalization (Fig. 4). The NCCL documentation[3] also details further the loading
 77 behavior:

- 78 • If `NCCL_PROFILER_PLUGIN` is set: attempt to load the library with the specified
 79 name;
 80 if that fails, attempt `libnccl-profiler-<NCCL_PROFILER_PLUGIN>.so`.
- 81 • If `NCCL_PROFILER_PLUGIN` is not set: attempt `libnccl-profiler.so`.
- 82 • If no plugin was found: profiling is disabled.
- 83 • If `NCCL_PROFILER_PLUGIN` is set to `STATIC_PLUGIN`, the plugin symbols are
 84 searched in the program binary.

85 The plugin loading mechanism expects the variable name to follow the naming convention
 86 `ncclProfiler_v{versionNum}`, which also indicates the API version.

87 The profiler API has changed multiple times with newer NCCL releases. NCCL features a fallback
 88 mechanism to load older versions. However one instance is known, where a profiler plugin being de-
 89 veloped against the NCCL release 2.25.1 with Profiler API version 2, was unable to run with the latest
 90 NCCL release[5]. Around this point in time, the NCCL repository has undergone a refactor related to
 91 the profiler plugin.

⁴src/plugin/profiler.cc.

⁵src/plugin/plugin_open.cc.



Figure 4: User API → NCCL communicator init → load profiler plugin and call `profiler->init()`.

92 3.2 Profiler API

93 The plugin must implement a profiler API specified by NCCL by exposing a struct⁶. This struct should
 94 contain pointers to all functions required by the API. A plugin may expose multiple versioned structs
 95 for backwards compatibility with older NCCL versions.

96

97 `ncclProfiler_v5_t ncclProfiler_v5 = {`

⁶src/include/plugin/profiler/profiler_v5.h.

```

98 const char* name;
99 ncclResult_t (*init)(...); // called when a communicator is created
100 ncclResult_t (*startEvent)(...); // at start of operations/activities
101 ncclResult_t (*stopEvent)(...); // at end of these operations/activities
102 ncclResult_t (*recordEventState)(...); // to record state of certain
103     operations
104 ncclResult_t (*finalize)(...); // called when a communicator is destroyed
105 };
```

- 107 As of NCCL v2.29.2, version 6 is the latest, which was released on Dec 24, 2025. This release
 108 happened well after the begin of the study, so the focus will remain on version 5. Version 6 introduced
 109 additional profiler features for Copy-Engine based collective operations, otherwise version 6 and version
 110 5 remain the same.
- 111 Five functions must be implemented for the API. Internally NCCL wraps all calls to the profiler API in
 112 custom functions which are all neatly declared in a single file⁷.
- 113 NCCL invokes the profiler API in different code regions to capture start and stop of NCCL groups,
 114 collectives, P2P, proxy, kernel and network activity. As the API function names suggest, this will allow
 115 the profiler to track these operations and activities as events.
- 116 The API functions and where NCCL invokes them are explained in the following sections.

117 3.2.1 init

- 118 `init` initializes the profiler plugin with a communicator. `init` is called immediately after
 119 `ncclProfilerPluginLoad`, which happens every time a communicator is created (see Fig. 4).
 120 This may happen multiple times for the same profiler instance, if further communicators are created on
 121 that process. NCCL passes following arguments:

```

122
123 ncclResult_t init(
124     void** context, // out param - opaque profiler context
125     uint64_t commId, // communicator id
126     int* eActivationMask, // out param - bitmask for which events are tracked
127     const char* commName, // user assigned communicator name
128     int nNodes, // number of nodes in communicator
129     int nranks, // number of ranks in communicator
130     int rank, // rank identifier in communicator
131     ncclDebugLogger_t logfn // logger function
132 );
133
```

- 134 If the profiler plugin `init` function does not return `ncclSuccess`, NCCL disables the plugin.
 135 `void** context` is an opaque handle that the plugin developer may point to any custom context

⁷[src/include/profiler.h](#).

136 object; this pointer is passed again in `startEvent` and `finalize`. This context object is separate
 137 per communicator.

138 The plugin developer should set `int* eActivationMask` to a bitmask, indicating which event
 139 types the profiler wants to track⁸:

```
140
141 enum {
142     ncclProfileGroup = (1 << 0), // group event type
143     ncclProfileColl = (1 << 1), // host collective call event type
144     ncclProfileP2p = (1 << 2), // host point-to-point call event type
145     ncclProfileProxyOp = (1 << 3), // proxy operation event type
146     ncclProfileProxyStep = (1 << 4), // proxy step event type
147     ncclProfileProxyCtrl = (1 << 5), // proxy control event type
148     ncclProfileKernelCh = (1 << 6), // kernel channel event type
149     ncclProfileNetPlugin = (1 << 7), // network plugin-defined, events
150     ncclProfileGroupApi = (1 << 8), // Group API events
151     ncclProfileCollApi = (1 << 9), // Collective API events
152     ncclProfileP2pApi = (1 << 10), // Point-to-Point API events
153     ncclProfileKernelLaunch = (1 << 11), // Kernel launch events
154 };
155
```

156 The default value is to 0, which means no events are tracked by the profiler. Setting it to 4095 will track
 157 all events.

158 `ncclDebugLogger_t logfn` is a function pointer to NCCL’s internal debug logger
 159 (`ncclDebugLog`). NCCL passes this so the plugin can emit log lines through the same channel and
 160 filtering as NCCL: the plugin may store the callback and call it with `(level, flags, file,`
 161 `line, fmt, ...)` when it wants to log. Messages then appear in NCCL’s debug output (e.g.
 162 `stderr` or `NCCL_DEBUG_FILE`) and respect the user’s `NCCL_DEBUG` level and subsystem mask. Using
 163 `logfn` keeps profiler output consistent with NCCL’s own logs.

164 3.2.2 startEvent

165 `startEvent` is called when NCCL begins certain operations:

```
166
167 ncclResult_t startEvent(
168     void* context, // opaque profiler context object
169     void** eHandle, // out param - event handle
170     ncclProfilerEventDescr_v5_t* eDescr // pointer to event descriptor
171 );
```

173 As of release v2.29.2 NCCL does not use the return value. `void** eHandle` may point to a cus-
 174 tom event object; this pointer is passed again in `stopEvent` and `recordEventState`. `eDescr`
 175 describes the started event⁹.

⁸[src/include/plugin/nccl_profiler.h](#)

⁹[src/include/plugin/profiler/profiler_v5.h](#)

176 The field `void* parentObj` in the event descriptor is the `eHandle` of a parent event (or null). The
 177 use of this field can be explained as following:

178 All User API calls to Collective or Point-to-Point operations will start a Group API event. When net-
 179 working is required, ProxyCtrl Events may be emitted. Depending on the `eActivationMask` bitmask
 180 returned in the `init` function, further (child) events will be emitted in deeper regions of the nccl code
 181 base. It can be thought of as an event hierarchy¹⁰ with several depth levels:

```

182
183     Group API event
184     |
185     +- Collective API event
186     |   |
187     |   +- Collective event
188     |   |
189     |   +- ProxyOp event
190     |   |   |
191     |   |   +- ProxyStep event
192     |   |   |
193     |   |   +- NetPlugin event
194     |   |
195     |   +- KernelCh event
196
197     +- Point-to-point API event
198     |   |
199     |   +- Point-to-point event
200     |   |
201     |   +- ProxyOp event
202     |   |   |
203     |   |   +- ProxyStep event
204     |   |   |
205     |   |   +- NetPlugin event
206     |   |
207     |   +- KernelCh event
208
209     +- Kernel Launch event
210
211     ProxyCtrl event
  
```

213 The `parentObj` inside `eDescr` will be a reference to the `eHandle` of the respective parent event
 214 for the current event according to this hierarchy. Thus, if the `eActivationMask` set during `init`
 215 enables tracking for event types lower in the hierarchy, NCCL always also tracks their parent event types.

216 3.2.3 stopEvent

```

217
218     ncclResult_t stopEvent (void* eHandle); // handle to event object
  
```

220 `stopEvent` tells the plugin that the event has stopped. `stopEvent` for collectives simply indicates
 221 to the profiler that the collective has been enqueued and not that the collective has been completed.

222 As of NCCL v2.29.2 NCCL does not use the return value.

223 `stopEvent` is called in the same functions that call `startEvent`, except for the `GroupApi` event.
 224 Fig. 5 shows when NCCL emits `startEvent` and `stopEvent` after a user API call. The Proxy-
 225 Progress thread also emits `startEvent` and `stopEvent` while progressing ops (see Fig. 6).

¹⁰`ext-profiler/README.md`.




Figure 5: Flow from NCCL API calls to profiler events. Internally, some Collectives (e.g. ncclAll-toAll) are implemented as multiple point-to-point operations, triggering many P2pApi and P2p events. For application settings utilizing ncclGroupStart/End, further examples and visualizations are provided in section 4.




Figure 6: `progressOps` emits `ProxyStep/KernelCh/NetPlugin` events. `getPostedOps` emits `ProxyOp` events. Several events `ProxyCtrl` are also emitted

226 `op->progress()` progresses transport specific ops. This is implemented as a function pointer type¹¹.
 227 Confusingly the variable is called ‘`op`’, although its type is `ncclProxyArgs` and *not* `ncclProxyOp`.
 228
 229 `typedef ncclResult_t (*proxyProgressFunc_t) (struct ncclProxyState*, struct`
 230 `ncclProxyArgs*);`
 231
 232 `struct ncclProxyArgs {`
 233 `proxyProgressFunc_t progress;`
 234 `struct ncclProxyArgs* next;`
 235 `/* other fields */`
 236 `}`
 237

238 This allows calls to different the implementations of the `progress` function for different transport

¹¹ `src/include/proxy.h`.

239 methods¹²¹³¹⁴¹⁵. Each implementations calls the profiler API to inform about a different event type
 240 (ProxyStep, KernelCh or Network plugin specific).

241 **3.2.4 recordEventState**

```
242 ncclResult_t recordEventState(  

243     void* eHandle,  

244     ncclProfilerEventState_v5_t eState,  

245     ncclProfilerEventStateArgs_v5_t* eStateArgs  

246 );  

247 
```

249 Some event types can be updated by NCCL through `recordEventState` (state and attributes)¹⁶.
 250 `recordEventState` is called in the same functions that call `startEvent`, while always being
 251 called after `startEvent`.

252 **3.2.5 finalize**

```
253 ncclResult_t finalize(void* context);  

254 
```

256 After a user API call to free resources associated with a communicator, `finalize` is called. Af-
 257 terwards, a reference counter tracks how many communicators are still being tracked by the profiler
 258 plugin. If it reaches 0, the plugin will be closed via `dlclose(handle)`. Fig. 7 depicts the flow from
 259 user API call to `finalize`.

¹²`src/transport/net.cc`.

¹³`src/transport/coll_net.cc`.

¹⁴`src/transport/p2p.cc`.

¹⁵`src/transport/shm.cc`.

¹⁶`src/include/plugin/profiler/profiler_v5.h`.




Figure 7: Flow from User API call → finalize and plugin unload.

260 3.2.6 name

261 The profiler plugin struct also has a `name` field. The `name` field should point to a character string with
 262 the name of the profiler plugin. It will be used for all logging, especially when `NCCL_DEBUG=INFO` is
 263 set.

264 4 Code examples and visualizations

265 The following examples illustrate the profiling behavior for different user application settings:

- 266 • One Device per Thread
- 267 • Multiple Devices per Thread via `ncclGroupStart` and `ncclGroupEnd`
- 268 • One Device per Thread and aggregated operations via `ncclGroupStart` and `ncclGroupEnd`

269 A profiler plugin that logs all call information to a file has been developed and is used in all examples.
 270 An exemplary illustration is shown below:

```

271
272 struct MyContext { /* custom context struct */ };
273 struct MyEvent { /* custom event struct */ };
274
275 MyEvent* allocEvent(args) { /* handles event allocation */ }

```

```
276 uint64_t getTime() { /* gets time */ }
277 void writeJsonl() { /* writes call details to process specific log file as
278     structured jsonl */ }
279
280 ncclResult_t myInit( /* args - **context, *eActivationMask, ... */ ) {
281     *context = malloc(sizeof(struct MyContext));
282     *eActivationMask = 4095; /* enable ALL event types */
283
284     writeJsonl(getTime(), "Init", args);
285     return ncclSuccess;
286 }
287
288 ncclResult_t myStartEvent( /* args - **eHandle, ... */ ) {
289     *eHandle = allocEvent(args);
290
291     writeJsonl(getTime(), "StartEvent", args);
292     return ncclSuccess;
293 }
294
295 ncclResult_t myStopEvent(void* eHandle) {
296     writeJsonl(getTime(), "StopEvent", eHandle);
297
298     free(eHandle)
299     return ncclSuccess;
300 }
301
302 ncclResult_t myRecordEventState( /* args - ... */ ) {
303     writeJsonl(getTime(), "RecordEventState", args);
304     return ncclSuccess;
305 }
306
307 ncclResult_t myFinalize(void* context) {
308     writeJsonl(getTime(), "Finalize", args);
309
310     free(context);
311     return ncclSuccess;
312 }
313
314 ncclProfiler_v5_t ncclProfiler_v5 = {
315     "MyProfilerPlugin",
316     myInit,
317     myStartEvent,
318     myStopEvent,
319     myRecordEventState,
320     myFinalize,
321 };
```

323 Alongside the logging profiler plugin, a visualization tool has been built, that ingests the profiler logs
 324 to inspect the exact behavior of internal calls from NCCL to the Profiler API. It displays the events
 325 as colored bars on a timeline and separates them on different lanes. Each lane also displays some
 326 information about the communicator, rank and thread corresponding to the event. Additionally, blue
 327 dotted lines indicate the relationship between events according to the `parentObj` field and red lines
 328 indicate which collective events belong to the same collective operation.
 329 Further, a hover feature was added to inspect all details of an event, however this feature is not used in
 330 the following illustrative examples.

331 4.1 One Device per Thread

332 This example visualizes an AllReduce collective across multiple GPUs (see Fig. 8 and Fig. 9). Each
 333 NCCL thread manages a single GPU. This may be achieved by starting out with the same number of
 334 MPI tasks with each task running single threaded; or by having less MPI tasks, but the tasks create
 335 multiple thread workers. Custom initialization without MPI is also possible if desired.

```
336
337 // broadcast a commId
338
339 // ...
340
341 ncclCommInitRank(&rootComm, nRanks, commId, myRank);
342
343 // ...
344
345 ncclAllReduce(sendBuff, recvBuff, BUFFER_SIZE, ncclFloat, ncclSum, rootComm
346     , streams);
347
348 // ...
349
350 ncclCommDestroy(rootComm);
```

352 The profiler API calls are visualized in Fig. 8 and Fig. 9. Below follows a full description of the calls to
 353 the profiler API induced by the example program:

354 First, the profiler API `init` is called for each rank. This occurs during NCCL's internal communicator
 355 creation, when the application calls `ncclCommInitRank`. After the application calls
 356 `ncclAllReduce`, many Profiler API calls to `startEvent`, `stopEvent`, and `recordEventState`
 357 are triggered: Initially, `startEvent` for the `groupApi` (green bar) is called. Below it, the `startEvent` and
 358 soon the `stopEvent` for the AllReduce `collApi` event are called. The yellow bar shows when NCCL
 359 enqueues the GPU kernel launch (`KernelLaunch` event). The two bars below represent the `group`
 360 and `coll` events. NCCL also spawns a proxy progress thread per rank, which does additional pro-
 361 filer API calls. The first red `ProxyCtrl` event shows the proxy progress thread was asleep. Next,
 362 a new `ProxyCtrl` event shows time for the proxy thread to append proxy ops. Then, appended ops
 363 start progressing triggering `ProxyOps` events, which in `op->progress()` starts `ProxyStep` and

364 KernelCh events that inform about low level network activity in updates via recordEventState
 365 like ProxyStepRecvGPUWait (see Fig. 9). Network activity eventually completes and the AllRe-
 366 duce collective finishes. The next ProxyCtrl event only shows the proxy thread sleeping again.
 367 Finally, profiler finalize is called, which happens when the application cleans up NCCL communi-
 368 cators and no further communicators are tracked in the profiler in each respective thread.

369 ProxyStep events are emitted in environments requiring cross node communication. If this type of
370 communication is not required, then ProxyStep events will not be emitted either.




Figure 8: One device per thread: A visualization of the calls generated to the Profiler API, starting from communicator creation, followed by a collective operation and communicator destruction. ProxyStep events have been omitted for visual clarity, see Fig. 9 for a depiction.




Figure 9: One device per thread: In Fig. 8 ProxyStep events have been omitted for visual clarity. However, in multinode settings, many additional profiler API calls for proxyStep events happen, informing about the low level network steps in their event details via recordEventState (indicated as red circles above each of the event bars). The blue dotted lines indicate the parentObj of each proxyStep event, which are the above proxyOp events.

371 4.2 Multiple Devices per Thread (ncclGroup)

372 In this example¹⁷, one NCCL thread manages all GPUs on the same node. This is achieved by wrapping
 373 communication initialization in ncclGroupStart and ncclGroupEnd for each managed GPU. In
 374 this orchestration setting, **NVIDIA’s documentation states that collective API calls should also be**
 375 **wrapped in ncclGroup**. Here, only one collective operation (per device) is inside the ncclGroup:

```

376
377 // broadcast a commId
378
379 // ...
380
  
```

¹⁷examples/03_collectives/01_allreduce/.

```
381 ncclGroupStart();  
382 for (int i=0; i<nngpus; i++) {  
383     cudaSetDevice(dev);  
384     ncclCommInitRank(comms+i, nngpus*nRanks, id, myRank*nngpus+i);  
385 }  
386 ncclGroupEnd();  
387  
388 // alternatively to above method, NCCL provides the convenience function  
389 // ncclCommInitAll();  
390  
391 // ...  
392  
393 ncclGroupStart();  
394 for (int i = 0; i < num_gpus; i++) {  
395     ncclAllReduce( /* ... */ );  
396 }  
397 ncclGroupEnd();  
398  
399 // ...  
400  
401 for (int i = 0; i < num_gpus; i++) {  
402     ncclCommDestroy(comms[i]);  
403 }
```

405 In this example case, the profiler API behavior remains largely the same: The one difference is that
406 NCCL internally calls the profiler API groupApi event only one time in total for aggregated operations
407 within a thread. Otherwise all other events are processed as usual and are called their usual amount of
408 times irrespective of ncclGroup (Fig. 10). This behaviour also holds true within a process. It also
409 holds when grouping (single) collectives for different communicators.




Figure 10: Multiple devices per thread: Events from the proxy thread as well as init and finalize calls are omitted. Collective API calls from multiple GPUs managed by a single thread only trigger a single GroupApi event.

4.3 Aggregated operations

In this example, the setting is such that only a single GPU is managed by a thread, but multiple collective operations are grouped (i.e. to optimize communication efficiency):

```

411 // broadcast a commId
412
413 // ...
414
415
416 ncclCommInitRank(&rootComm, nRanks, rootId, myRank);
417
418 // ...
419
420 ncclGroupStart();
421
422 ncclAllReduce( /* ... */, rootComm, /* ... */ );
423 ncclBroadcast( /* ... */, rootComm, /* ... */ );
424 ncclReduce( /* ... */, rootComm, /* ... */ );
425 ncclAllGather( /* ... */, rootComm, /* ... */ );
426 ncclReduceScatter( /* ... */, rootComm, /* ... */ );
427 ncclGroupEnd();
428
429 // ...
430
431

```

The behavior changes can be described as follow:

- 433 • single GroupApi event per thread
 434 • single KernelLaunch event per thread
 435 • single Group event per thread




Figure 11: one GPU per thread with aggregated operations: multiple collective calls are grouped together and nccl does only a single kernel launch per thread.

436 5 Performance and scalability

437 Experiments were run to assess the performance and scalability of profiler plugins. These experiments
 438 measure the overhead induced internally by NCCL to serve the profiler plugin, but do not intend to
 439 measure the performance of a profiler plugin itself as the plugin is fully customizable to the needs of the
 440 developer.

441 Thus, the profiler developed for the experiments only initializes a dummy context struct, returns NULL
 442 for event handles and tracks all events (eActivationMask set to 4095).

```
443
444 // an 'empty' NCCL Profiler Plugin
445
446 struct MyContext {
447     char dummy;
448 };
449
450 ncclResult_t myInit(void** context, uint64_t commId, int* eActivationMask,
```

```

451     const char* commName, int nNodes, int nranks, int rank,
452     ncclDebugLogger_t logfn) {
453     *context = malloc(sizeof(struct MyContext));
454     *eActivationMask = 4095; /* enable ALL event types */
455     return ncclSuccess;
456 }
457
458 ncclResult_t myStartEvent(void* context, void** eHandle,
459     ncclProfilerEventDescr_v5_t* eDescr) {
460     *eHandle = NULL;
461     return ncclSuccess;
462 }
463
464 ncclResult_t myStopEvent(void* eHandle) {
465     return ncclSuccess;
466 }
467
468 ncclResult_t myRecordEventState(void* eHandle, ncclProfilerEventState_v5_t
469     eState, ncclProfilerEventStateArgs_v5_t* eStateArgs) {
470     return ncclSuccess;
471 }
472
473 ncclResult_t myFinalize(void* context) {
474     free(context);
475     return ncclSuccess;
476 }
477
478 ncclProfiler_v5_t ncclProfiler_v5 = {
479     "EmptyProfiler",
480     myInit,
481     myStartEvent,
482     myStopEvent,
483     myRecordEventState,
484     myFinalize,
485 };

```

487 For testing the performance overhead in collective and P2P operations, **nccl-tests** from NVIDIA was
 488 used¹⁸.

489 The applications `sendrecv_perf` and `all_reduce_perf` were launched with following test pa-
 490 rameters: message size 64 B, 1 000 000 iterations per size, 100 warmup iterations. Single-node jobs used
 491 one node and 4 GPUs; multi-node jobs used 2 nodes, 4 GPUs per node, 8 MPI ranks in total. For each
 492 experiment, the application was run once without the profiler and once with the empty profiler plugin.

493 All experiments were carried out on nodes of the ZIH HPC Capella cluster. Each node is equipped with
 494 4 x Nvidia H100 GPUs and 2 x AMD EPYC 9334 (32 cores) @ 2.7 GHz CPUs.

¹⁸<https://github.com/NVIDIA/nccl-tests>

495 Table 1 shows the average latency per operation (time in μ s) across iterations. The empty profiler adds
 496 roughly 8 μ s to 9 μ s overhead per operation in single-node runs (4 GPUs), but introduces negligible
 497 overhead in multi-node runs (8 GPUs across 2 nodes).

Table 1: Profiler overhead: nccl-tests sendrecv_perf (P2P) and all_reduce_perf (collectives).

Latency averaged over 1M iterations.

| Test | Environment | Without profiler (μ s) | With profiler (μ s) |
|-------------------------------|-----------------------|-----------------------------|--------------------------|
| P2P (sendrecv_perf) | Single-node (4 GPUs) | 14.3 | 23.88 |
| | Multi-node (2×4 GPUs) | 13.05 | 12.95 |
| Collectives (all_reduce_perf) | Single-node (4 GPUs) | 14.96 | 23.29 |
| | Multi-node (2×4 GPUs) | 17.99 | 18.34 |

498 Using the profiler plugin when scaled to many gpus across multiple nodes is effortless and did not require
 499 any changes to the profiler plugin for the used code examples and experiments.

500 6 Discussion

501 This section first discusses practical considerations for developers who implement or extend a NCCL
 502 profiler plugin, as well as known limitations of the current profiling infrastructure, and then shows how
 503 the plugin could be integrated with the Score-P measurement infrastructure for HPC-wide tracing and
 504 analysis.

505 6.1 Considerations for developers

506 **Profiler visualization.** The visualization tool used in the code examples is helpful for understanding
 507 the internal call behavior to the Profiler API by NCCL and will be made available along with this
 508 report. It may serve as a reference to compare against for other developers that build a profiler plugin or
 509 visualizer

510 **Correlating Collective Events with seqNumber.** When profiling is enabled, NCCL counts the
 511 number of calls for each type of collective function per communicator.

```
512 /src/include/comm.h
513
514 struct ncclComm {
515     uint64_t seqNumber[NCCL_NUM_FUNCTIONS];
516     /* other fields */
517 }
```

519 /src/plugin/profiler.cc

```
520
521 ncclResult_t ncclProfilerStartTaskEvents(struct ncclKernelPlan* plan) {
522     /* other code */
523     __atomic_fetch_add(&plan->comm->seqNumber[ct->func], 1, __ATOMIC_RELAXED)
524     ;
525     /* other code */
526 }
```

528 This value is present in the `eDescr` for collective events and can be used to identify which collectives
529 operations belong together across processes (see Fig. 12).

530 **Tracing low level activity back to NCCL API calls with parentObj.** If a plugin developer
531 wants utilize this field, they should ensure that potential address reuse does not create ambiguity to
532 what the parentObj was originally pointing to. *Custom memory management is advised.* This field is
533 useful when trying to understand which user API call triggered which events of lower level operations
534 or activity such as network activity (see Fig. 12).




Figure 12: An example illustrating how `parentObj` and `seqNumber` can be used to better understand the timing of concurrent collective operations.

535 **Process origin for profiler callbacks with PXN enabled.** Unless Setting the environment vari-
536 able NCCL_PXN_DISABLE=0 (default 1), due to PXN (PCIe x NVLink) some proxy ops may be
537 progressed in a proxy thread from another process, different to the one that originally generated the
538 operation. Then parentObj in eDescr is not safe to dereference; the eDescr for ProxyOp events
539 includes the originator's PID, which the profiler can match against the local PID. The eDescr for
540 ProxyStep does not provide this field. However a workaround is possible:

541 The passed context object in startEvent is also unsafe to dereference due to PXN. the profiler
542 plugin developer may internally track initialized contexts and whether the passed context belongs to

543 the local process. This is also indicative of PNX.

544 **Tracking communicator parent–child relationships.** With the current Profiler plugin API, it is
545 not possible to detect whether a communicator originates from another one (e.g., via
546 `ncclCommSplit` or `ncclCommShrink`). The plugin’s `init` callback only receives a single com-
547 munication ID (`commId`, which corresponds to `comm->commHash`), as well as `commName`, `nNodes`,
548 `nRanks`, and `rank`; there is no `parentCommId` or similar argument. In split/shrink, the `commHash`
549 of the child node is calculated internally as a one-way digest of the `commHash` of the parent node and
550 the split parameters (`splitCount`, `color`). Therefore, the relationship cannot be restored based on
551 the ID alone.

552 6.2 Known limitations

553 Kernel event instrumentation uses counters exposed by the kernel to the host and the proxy progress
554 thread. Thus the proxy progress thread infrastructure is shared between network and profiler. If the
555 proxy is serving network requests, reading kernel profiling data can be delayed, causing loss of accuracy.
556 Similarly, under heavy CPU load and delayed scheduling of the proxy progress thread, accuracy can be
557 lost.

558 From profiler version 4, NCCL uses a per-channel ring buffer of 64 elements. Each counter is comple-
559 mented by two timestamps (ptimers) supplied by the NCCL kernel (start and stop of the operation in
560 the kernel). NCCL propagates these timestamps to the profiler plugin so it can convert them to the CPU
561 time domain.¹⁹

562 6.3 Potential Integration with Score-P

563 The Score-P measurement infrastructure[1] is a highly scalable and easy-to-use tool suite for profiling
564 and event tracing of HPC applications. It supports a number of analysis tools. Currently, it works
565 with Scalasca, Vampir, and Tau and is open for other tools and produces OTF2 traces and CUBE4
566 profiles. Integrating NCCL into Score-P allows developers to see communication collectives alongside
567 the application logic.

568 A prerequisite for distributed tracing is the unique identification of process groups. NCCL achieves this
569 via `ncclGetUniqueId`²⁰ without a central coordinator. To establish a communicator, one process
570 generates a handle containing a random 64-bit magic value from `/dev/urandom` and the socket
571 address of a new listening socket (IP, port), whose port is chosen by the operating system. This
572 combination avoids collisions across a cluster, ensuring the ID is globally unique in practice. A Score-p
573 integration can use these to accurately define Process Groups.

¹⁹`ext-profiler/README.md`.

²⁰`src/init.cc`.

- 574 The integration could be achieved in two ways, either using a direct Profiler API mapping or via an
 575 indirect NVTX/CUPTI annotation:
- 576 A direct integration would potentially involve implementing a NCCL profiler plugin that translates the
 577 startEvent and stopEvent callbacks into Score-P regions: The plugin maps NCCL event descrip-
 578 tors (e.g., ncclAllReduce) to Score-P regions using Score-P instrumentation (e.g.,
 579 SCOREP_EnterRegion and SCOREP_ExitRegion or
 580 SCOREP_USER_REGION_BY_NAME_BEGIN/END).
- 581 Alternatively, the NCCL profiler plugin can act as a bridge to NVIDIA’s Tools Extension (NVTX). If
 582 Score-P has been built with CUDA support it can intercept NVTX ranges. The NCCL profiler plugin
 583 would emit nvtxRangePush[4] and nvtxRangePop around NCCL operations. Score-P records
 584 these as labeled regions without requiring the plugin to link directly against Score-P libraries. This ap-
 585 proach decouples the NCCL plugin from the Score-P build environment and instead relies on Score-P’s
 586 internal NVTX-to-OTF2 mapping logic.
- 587 The plugin could also directly utilize cuptiActivityPush/PopExternalCorrelationId to
 588 capture GPU activity during the startEvent and stopEvent of KernelLaunch events, while
 589 incrementing a thread-safe correlation ID (see Fig. 13). CUPTI can be initialized and cleaned up within
 590 the profiler plugin’s init and finalize functions.




Figure 13: CUPTI activity is visualized as orange event bars. With a unique correlation Id, it is possible to trace the activity back to KernelLaunch events

591 7 Conclusion

- 592 This study examined the NCCL Profiler Plugin API and its suitability for integration with Score-P. It
 593 provided background on NCCL and its design, explained how the profiler plugin is loaded and described

594 the API definition with its five core callbacks `init`, `startEvent`, `stopEvent`,
595 `recordEventState` and `finalize`. Code examples and visualizations illustrate the event flow
596 from API calls to NCCL’s internal profiler callbacks. Performance experiments showed that an empty
597 profiler adds roughly 8 µs to 9 µs overhead per operation in single-node runs but introduces negligible
598 overhead in multi-node runs, and scaling to many GPUs across nodes required no changes to the profiler
599 plugin. The discussion covered developer considerations, known limitations, and a potential integration
600 strategy with Score-P.

601 The NCCL Profiler API allows for highly customized plugins tailored to the analysis needs, whether for
602 simple timing, kernel tracing via CUPTI, or integration with external tools such as Score-P. A notable
603 advantage is its low overhead: NVIDIA advertises their `inspector`²¹ implementation as efficient
604 enough for “always-on” profiling in production. On the downside, profiler plugins may require mainte-
605 nance and active development, since NCCL is actively developed. API versions evolve and new features
606 are being introduced.

²¹`ext-profiler/inspector`.

607 References

608 [1] Andreas Knüpfer et al. Score-p: Performance measurement infrastructure. <https://www.vi-hps.org/projects/score-p/overview>.

610 [2] NVIDIA. NCCL: Nvidia collective communication library. <https://github.com/NVIDIA/nccl>.

612 [3] NVIDIA. Nccl user guide. <https://docs.nvidia.com/deeplearning/nccl/user-guide/docs/env.html#nccl-profiler-plugin>.

614 [4] NVIDIA. Nvtx: Markers and ranges. https://nvidia.github.io/NVTX/doxygen/group__m_a_r_k_e_r_s__a_n_d__r_a_n_g_e_s.html.

616 [5] Ioannis Vardas. ncclsee. <https://github.com/variemai/ncclsee>.